

## Project 2- Regression

### Aadith Narayan Ravishankar AXR180085

Source: random sample of 1734 houses taken from full Saratoga Housing Data (De Veaux) Number of Cases: 1734 Story: House prices and expenses in New York. What properties of a house can predict its price? Can we use such a model to identify houses that are extraordinarily expensive or inexpensive?

Data Description.

- price price(1000s of US dollars)
- lotSize size of lot (square feet)
- age age of house (years)
- landValue value of land (1000s of US dollars)
- livingArea living area (square feet)
- pctCollege percent of neighborhood that graduated college
- bedrooms number of bedrooms
- fireplaces number of fireplaces
- batrooms number of batrooms (half batrooms have no shower or tub)
- rooms number of rooms
- heatingType type of heating system
- fuelType fuel used for heating system
- sewerType type of sewer system
- waterfront whether property includes waterfront
- newConstruct whether the property is a new construction
- centralAir whether the house has central air

## Pre-Processing of the dataset

### Exploratory Data Analysis

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]:
```

```
In [5]: house=pd.read_csv("house.csv")
```

```
In [6]: house.head()
```

```
Out [6]:
```

|   | Price  | LotSize | Waterfront | Age  | LandValue | NewConstruct | CentralAir | FuelType | HeatType  | SewerType | LivingArea | PctCollege |
|---|--------|---------|------------|------|-----------|--------------|------------|----------|-----------|-----------|------------|------------|
| 0 | 317000 | 0.46    | 0          | 17.0 | 60500     | 0            | 1          | Gas      | Hot Water | Public    | 1762       | 63.0       |
| 1 | 182000 | 0.30    | 0          | 32.0 | 22500     | 0            | 0          | Electric | Electric  | Private   | 1753       | 35.0       |
| 2 | 297000 | 0.30    | 0          | 15.0 | 34000     | 0            | 1          | Gas      | Hot Air   | Public    | 2288       | 84.0       |
| 3 | 210000 | 0.49    | 0          | 15.0 | 25000     | 0            | 0          | Gas      | Hot Air   | NaN       | 1592       | 54.0       |
| 4 | 130000 | 0.16    | 0          | 36.0 | 54000     | 0            | 0          | Gas      | Hot Air   | Public    | 1134       | 57.0       |

```
In [7]: df=pd.DataFrame(house)
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1734 entries, 0 to 1733
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  --
0   Price       1734 non-null    int64
1   LotSize     1734 non-null    float64
2   Waterfront  1734 non-null    int64
3   Age         1734 non-null    int64
4   LandValue   1734 non-null    int64
5   NewConstruct 1734 non-null    int64
6   CentralAir  1734 non-null    int64
7   FuelType    1734 non-null    object
8   HeatType    1734 non-null    object
9   SewerType   1709 non-null    object
10  LivingArea  1734 non-null    int64
11  PctCollege  1712 non-null    float64
12  Bedrooms   1734 non-null    int64
13  Fireplaces  1724 non-null    float64
14  Bathrooms  1734 non-null    float64
15  Rooms      1720 non-null    float64
16  Test       1734 non-null    int64
dtypes: float64(6), int64(8), object(3)
memory usage: 230.4+ KB
```

```
In [9]: df.isnull().sum()
```

```
Out [9]:
```

|       | Price         | LotSize     | Waterfront  | Age         | LandValue     | NewConstruct | CentralAir  | LivingArea  | PctCollege  | B   |
|-------|---------------|-------------|-------------|-------------|---------------|--------------|-------------|-------------|-------------|-----|
| count | 1734.000000   | 1734.000000 | 1734.000000 | 1734.000000 | 1734.000000   | 1734.000000  | 1734.000000 | 1734.000000 | 1712.000000 | 173 |
| mean  | 211545.064210 | 0.500294    | 0.006651    | 28.201985   | 34536.228374  | 0.046713     | 0.366205    | 1752.630911 | 55.639019   |     |
| std   | 98563.809881  | 0.698201    | 0.002632    | 29.887785   | 34980.940615  | 0.211084     | 0.481905    | 620.224953  | 10.303985   |     |
| min   | 5000.000000   | 0.000000    | 0.000000    | 0.000000    | 200.000000    | 0.000000     | 0.000000    | 616.000000  | 20.000000   |     |
| 25%   | 145000.000000 | 0.170000    | 0.000000    | 13.000000   | 15100.000000  | 0.000000     | 0.000000    | 1300.000000 | 52.000000   |     |
| 50%   | 189700.000000 | 0.370000    | 0.000000    | 19.000000   | 25000.000000  | 0.000000     | 0.000000    | 1632.000000 | 57.000000   |     |
| 75%   | 257289.000000 | 0.540000    | 0.000000    | 34.000000   | 40200.000000  | 0.000000     | 1.000000    | 2133.500000 | 64.000000   |     |
| max   | 775000.000000 | 12.200000   | 1.000000    | 225.000000  | 412600.000000 | 1.000000     | 1.000000    | 5228.000000 | 82.000000   |     |

```
In [11]: df.drop(['Test'],axis=1,inplace=True)
```

Imputing Null values in the dataset

```
In [12]: df['Age']=df[['Age']].transform(lambda x: x.fillna(int(x.median())))
```

```
In [13]: df['SewerType'].unique()
```

```
Out [13]: array(['Public', 'Private', nan, 'None/Unknown'], dtype=object)
```

```
In [14]: df['SewerType']=df[['SewerType']].transform(lambda x: x.fillna('None/Unknown'))
```

```
In [15]: df.describe()
```

```
Out [15]:
```

|       | Price         | LotSize     | Waterfront  | Age         | LandValue     | NewConstruct | CentralAir  | LivingArea  | PctCollege  | B   |
|-------|---------------|-------------|-------------|-------------|---------------|--------------|-------------|-------------|-------------|-----|
| count | 1734.000000   | 1734.000000 | 1734.000000 | 1734.000000 | 1734.000000   | 1734.000000  | 1734.000000 | 1734.000000 | 1712.000000 | 173 |
| mean  | 211545.064210 | 0.500294    | 0.006651    | 28.201985   | 34536.228374  | 0.046713     | 0.366205    | 1752.630911 | 55.639019   |     |
| std   | 98563.809881  | 0.698201    | 0.002632    | 29.887785   | 34980.940615  | 0.211084     | 0.481905    | 620.224953  | 10.303985   |     |
| min   | 5000.000000   | 0.000000    | 0.000000    | 0.000000    | 200.000000    | 0.000000     | 0.000000    | 616.000000  | 20.000000   |     |
| 25%   | 145000.000000 | 0.170000    | 0.000000    | 13.000000   | 15100.000000  | 0.000000     | 0.000000    | 1300.000000 | 52.000000   |     |
| 50%   | 189700.000000 | 0.370000    | 0.000000    | 19.000000   | 25000.000000  | 0.000000     | 0.000000    | 1632.000000 | 57.000000   |     |
| 75%   | 257289.000000 | 0.540000    | 0.000000    | 34.000000   | 40200.000000  | 0.000000     | 1.000000    | 2133.500000 | 64.000000   |     |
| max   | 775000.000000 | 12.200000   | 1.000000    | 225.000000  | 412600.000000 | 1.000000     | 1.000000    | 5228.000000 | 82.000000   |     |

```
In [16]: df['PctCollege']=df[['PctCollege']].transform(lambda x: x.fillna(int(x.median())))
```

```
In [17]: df['Fireplaces']=df[['Fireplaces']].transform(lambda x: x.fillna(int(x.median())))
```

```
In [18]: df['Rooms']=df[['Rooms']].transform(lambda x: x.fillna(x.median()))
```

```
In [19]: df.isnull().sum()
```

```
Out [19]:
```

|       | Price         | LotSize     | Waterfront  | Age         | LandValue     | NewConstruct | CentralAir  | LivingArea  | PctCollege  | B   |
|-------|---------------|-------------|-------------|-------------|---------------|--------------|-------------|-------------|-------------|-----|
| count | 1734.000000   | 1734.000000 | 1734.000000 | 1734.000000 | 1734.000000   | 1734.000000  | 1734.000000 | 1734.000000 | 1712.000000 | 173 |
| mean  | 211545.064210 | 0.500294    | 0.006651    | 28.201985   | 34536.228374  | 0.046713     | 0.366205    | 1752.630911 | 55.639019   |     |
| std   | 98563.809881  | 0.698201    | 0.002632    | 29.887785   | 34980.940615  | 0.211084     | 0.481905    | 620.224953  | 10.303985   |     |
| min   | 5000.000000   | 0.000000    | 0.000000    | 0.000000    | 200.000000    | 0.000000     | 0.000000    | 616.000000  | 20.000000   |     |
| 25%   | 145000.000000 | 0.170000    | 0.000000    | 13.000000   | 15100.000000  | 0.000000     | 0.000000    | 1300.000000 | 52.000000   |     |
| 50%   | 189700.000000 | 0.370000    | 0.000000    | 19.000000   | 25000.000000  | 0.000000     | 0.000000    | 1632.000000 | 57.000000   |     |
| 75%   | 257289.000000 | 0.540000    | 0.000000    | 34.000000   | 40200.000000  | 0.000000     | 1.000000    | 2133.500000 | 64.000000   |     |
| max   | 775000.000000 | 12.200000   | 1.000000    | 225.000000  | 412600.000000 | 1.000000     | 1.000000    | 5228.000000 | 82.000000   |     |

```
In [20]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1734 entries, 0 to 1733
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  --
0   Price       1734 non-null    int64
1   LotSize     1734 non-null    float64
2   Waterfront  1734 non-null    int64
3   Age         1734 non-null    int64
4   LandValue   1734 non-null    int64
5   NewConstruct 1734 non-null    int64
6   CentralAir  1734 non-null    int64
7   FuelType    1734 non-null    object
8   HeatType    1734 non-null    object
9   SewerType   1709 non-null    object
10  LivingArea  1734 non-null    int64
11  PctCollege  1734 non-null    float64
12  Bedrooms   1734 non-null    int64
13  Fireplaces  1734 non-null    float64
14  Bathrooms  1734 non-null    float64
15  Rooms      1734 non-null    float64
dtypes: float64(6), int64(7), object(3)
memory usage: 216.4+ KB
```

```
In [21]: df['FuelType'].unique()
```

```
Out [21]: array(['Gas', 'Electric', 'Oil', 'Unknown/Other', 'Wood', 'None', 'Solar'], dtype=object)
```

### Using one hot encoding to convert categorical features to numerical features

```
In [22]: cl=pd.get_dummies(df[['FuelType']],columns='FuelType',prefix='Fuel')
df[cl.columns]= cl
df.drop(['FuelType'],axis=1, inplace=True)
```

```
In [23]: df['HeatType'].unique()
```

```
Out [23]: array(['Hot Water', 'Electric', 'Hot Air', 'None'], dtype=object)
```

```
In [24]: cl=pd.get_dummies(df[['HeatType']],columns='HeatType',prefix='Heat')
df[cl.columns]= cl
df.drop(['HeatType'],axis=1, inplace=True)
```

```
In [25]: cl=pd.get_dummies(df[['SewerType']],columns='SewerType',prefix='Sewer')
df[cl.columns]= cl
df.drop(['SewerType'],axis=1, inplace=True)
```

```
In [26]: df.shape
```

```
Out [26]: (1734, 27)
```

### Splitting the data into train and test sets

```
In [27]: X=df.drop(['Price'],axis=1)
Y=df['Price']
```

```
In [28]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
```

```
In [29]: X_train_org,X_test_org,Y_train,Y_test= train_test_split(X,Y,test_size=0.2, random_state=0)
```

```
In [30]: scaler=MinMaxScaler()
scaler.fit(X_train_org)
X_train= scaler.transform(X_train_org)
X_test= scaler.transform(X_test_org)
print(X_train.shape)
(1387, 26)
```

### Bagging Regressor with Decision Tree Regressor

```
In [31]: from sklearn.ensemble import BaggingRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV

dt_clf = DecisionTreeRegressor(random_state=0)
bag_clf = BaggingRegressor(dt_clf, bootstrap=True, n_jobs=-1, random_state=0)
param_grid={'n_estimators':[50,100,150], 'max_samples':[100,200,300]}
grid=GridSearchCV(bag_clf,param_grid,cv=5,return_train_score=True)
grid.fit(X_train, Y_train)
train=grid.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid.score(X_test,Y_test)))

Best Parameters: {'max_samples': 300, 'n_estimators': 100}
Train score: 0.7487440603707227
Test score: 0.660601099744523
```

### Bagging Regressor with Linear SVR

```
In [32]: from sklearn.svm import SVR, LinearSVR
clf1=LinearSVR()
bag_clf1 = BaggingRegressor(clf1, bootstrap=True, n_jobs=-1, random_state=0)
param_grid={'n_estimators':[50,100,150], 'learning_rate':[0.5,1]}
grid=GridSearchCV(bag_clf1,param_grid,cv=5,return_train_score=True)
grid.fit(X_train, Y_train)
train=grid.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid.score(X_test,Y_test)))

Best Parameters: {'max_samples': 300, 'n_estimators': 50}
Train score: -5.41799945354148
Test score: -5.514295947942502
```

### Pasting with Linear Regression

```
In [33]: from sklearn.linear_model import LinearRegression
lr_clf = LinearRegression()
pas_clf = BaggingRegressor(lr_clf,bootstrap=False, n_jobs=-1, random_state=0)
param_grid={'n_estimators':[50,100,150], 'max_samples':[100,200,300]}
grid=GridSearchCV(pas_clf,param_grid,cv=5,return_train_score=True)
grid2.fit(X_train, Y_train)
train=grid2.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid2.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid2.score(X_test,Y_test)))

Best Parameters: {'max_samples': 300, 'n_estimators': 50}
Train score: 0.6531192708568357
Test score: 0.6558298304823773
```

### Pasting with Ridge Regression

```
In [34]: from sklearn.linear_model import Ridge
lr_clf = Ridge(alpha=0.1)
pas_clf1 = BaggingRegressor(lr_clf,bootstrap=False, n_jobs=-1, random_state=0)
param_grid={'n_estimators':[50,100,150], 'max_samples':[100,200,300]}
grid=GridSearchCV(pas_clf1,param_grid,cv=5,return_train_score=True)
grid2.fit(X_train, Y_train)
train=grid2.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid2.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid2.score(X_test,Y_test)))

Best Parameters: {'max_samples': 300, 'n_estimators': 50}
Train score: 0.658847479954797
Test score: 0.608686018227951
```

### Ada Boosting with Decision Tree Regressor

```
In [35]: from sklearn.ensemble import AdaBoostRegressor
ada_clf = AdaBoostRegressor(DecisionTreeRegressor(max_depth=1), random_state=0)
param_grid={'n_estimators':[50,100,150], 'learning_rate':[0.5,1]}
grid=GridSearchCV(ada_clf,param_grid,cv=5,return_train_score=True)
grid1.fit(X_train, Y_train)
train=grid1.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid1.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid1.score(X_test,Y_test)))

Best Parameters: {'learning_rate': 0.5, 'n_estimators': 50}
Train score: 0.24456061932372
Test score: 0.0904144237944822
```

### Ada Boosting with Lasso Regression

```
In [36]: from sklearn.linear_model import Lasso
ada_clf1 = AdaBoostRegressor(Lasso(alpha=10), random_state=0)
param_grid={'n_estimators':[50,100,150], 'learning_rate':[0.5,1]}
grid=GridSearchCV(ada_clf1,param_grid,cv=5,return_train_score=True)
grid5.fit(X_train, Y_train)
train=grid5.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid5.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid5.score(X_test,Y_test)))

Best Parameters: {'learning_rate': 1, 'n_estimators': 50}
Train score: 0.6090201881608291
Test score: 0.6190785936891614
```

### Gradient Boosting Regressor

```
In [37]: from sklearn.ensemble import GradientBoostingRegressor
gbt = GradientBoostingRegressor(max_depth=2, random_state=0)
param_grid={'n_estimators':[50,100,150], 'learning_rate':[0.5,1]}
grid=GridSearchCV(gbt,param_grid,cv=5,return_train_score=True)
grid6.fit(X_train, Y_train)
train=grid6.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid6.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid6.score(X_test,Y_test)))

Best Parameters: {'learning_rate': 0.5, 'n_estimators': 50}
Train score: 0.8937514177117686
Test score: 0.5921139788581473
```

### PCA with KNN Regression

```
In [38]: from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsRegressor
from sklearn.pipeline import Pipeline

clf = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("knn", KNeighborsRegressor())
])
param_grid={'n_neighbors': [3,5,7,9,19], 'knn_weights': ['uniform', 'distance'], 'knn_metric': ['l1', 'manhattan']}
grid7=GridSearchCV(clf,param_grid,cv=3,return_train_score=True)
grid7.fit(X_train,Y_train)
train=grid7.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid7.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid7.score(X_test,Y_test)))

Best Parameters: {'knn_metric': 'manhattan', 'knn_n_neighbors': 19, 'knn_weights': 'distance'}
Train score: 0.7901928020280424
Test score: 0.5115595216235551
```

### PCA with Linear Regression

```
In [39]: clf1 = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("lr", LinearRegression())
])
parameters=({'lr_fit_intercept':[True,False], 'lr_normalize':[True,False], 'lr_copy_X':[True, False]})
grid8=GridSearchCV(clf1,parameters,cv=3,return_train_score=True)
grid8.fit(X_train,Y_train)
train=grid8.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid8.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid8.score(X_test,Y_test)))

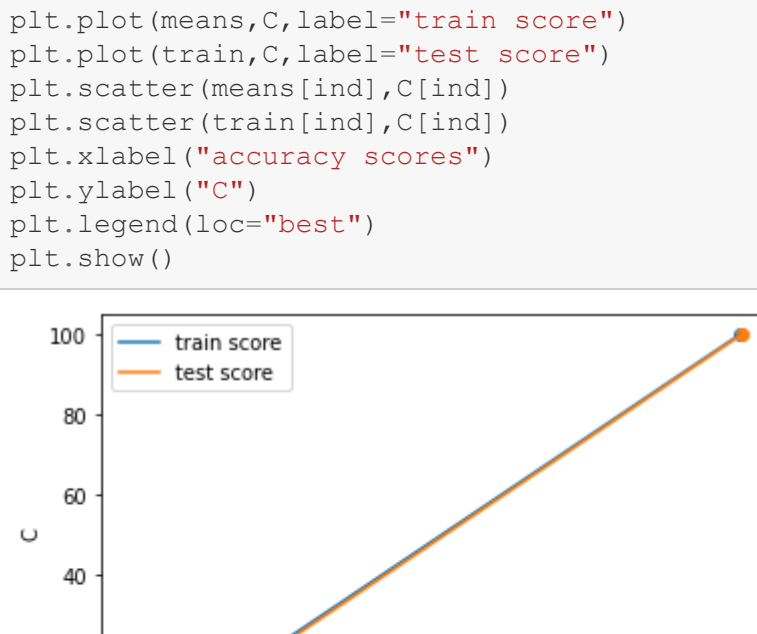
Best Parameters: {'lr_copy_X': True, 'lr_fit_intercept': True, 'lr_normalize': False}
Train score: -1.77331920117229
Test score: 0.4783450608838277
```

### PCA with Ridge Regression

```
In [40]: clf2 = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("rid", Ridge())
])
parameters1=({'rid_alpha':[0.1,1,10,100]}
grid9=GridSearchCV(clf2,parameters,cv=3,return_train_score=True)
grid9.fit(X_train,Y_train)
train=grid9.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid9.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid9.score(X_test,Y_test)))

Best Parameters: {'rid_alpha': 1}
Train score: 0.4297418518317177
Test score: 0.479003887649017
```

```
In [41]: alpha = [0.1,1,10,100]
means = grid9.cv_results_['mean_test_score']
for i in alpha:
    ind=grid9.best_params_['rid_alpha']:
        ind=alpha.index(i)
plt.plot(means,alpha,label="train score")
plt.plot(train,alpha,label="test score")
plt.scatter(means[ind],alpha[ind])
plt.scatter(train[ind],alpha[ind])
plt.xlabel("accuracy score")
plt.ylabel("alpha")
plt.legend(loc="upper right")
plt.show()
```

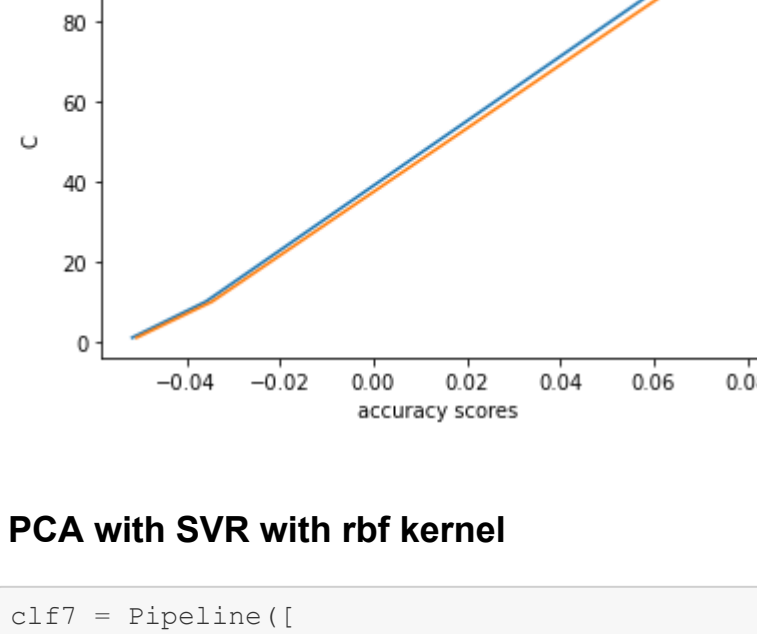


### PCA with Lasso Regression

```
In [42]: clf3 = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("las", Lasso())
])
parameters2=({'las_alpha':[0.1,1,10,100]}
grid10=GridSearchCV(clf3,parameters,cv=3,return_train_score=True)
grid10.fit(X_train,Y_train)
train=grid10.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid10.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid10.score(X_test,Y_test)))

Best Parameters: {'las_alpha': 100}
Train score: 0.4529353576760786
Test score: 0.47961120959591877
```

```
In [43]: alpha = [0.1,1,10,100]
means = grid10.cv_results_['mean_test_score']
for i in alpha:
    ind=grid10.best_params_['las_alpha']:
        ind=alpha.index(i)
plt.plot(means,alpha,label="train score")
plt.plot(train,alpha,label="test score")
plt.scatter(means[ind],alpha[ind])
plt.scatter(train[ind],alpha[ind])
plt.xlabel("accuracy score")
plt.ylabel("alpha")
plt.legend(loc="best")
plt.show()
```



### PCA with Polynomial Regression

```
In [44]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
pipe = make_pipeline(
    PolynomialFeatures(),
    Ridge()
)
clf4 = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("poly", pipe)
])
parameters2=({'poly_polynomialfeatures_degree':[1,2,3], 'poly_ridge_alpha':[0.1,1,10,100]}
grid11=GridSearchCV(clf4,parameters,cv=3,return_train_score=True)
grid11.fit(X_train,Y_train)
train=grid11.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid11.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid11.score(X_test,Y_test)))

Best Parameters: {'poly_polynomialfeatures_degree': 2, 'poly_ridge_alpha': 0.1}
Train score: 0.4801928020280424
Test score: 0.5345631769573334
```

### PCA with Linear SVR

```
In [45]: clf5 = Pipeline([
    ("pca", PCA(n_components=0.95)),
    ("lsvr", LinearSVR())
])
parameters2=({'lsvr_C':[1,10,100]}
grid12=GridSearchCV(clf5,parameters2,cv=3,return_train_score=True)
grid12.fit(X_train,Y_train)
train=grid12.cv_results_['mean_train_score']
print("Best Parameters: {}".format(grid12.best_params_))
print("Train score: {}".format(train.mean()))
print("Test score: {}".format(grid12.score(X_test,Y_test)))

Best Parameters: {'lsvr_C': 100}
Train score: -3.3461033537977
Test score: -0.991896257558669
```

```
In [46]: C=[1,10,100]
means = grid12.cv_results_['mean_test_score']
for i in C:
    if i==grid12.best_params_['lsvr_C']:
        ind=C.index(i)
plt.plot(means,C,label="train score")
plt.plot(train
```



```
[56]: # Step 1: Build the model
      modell = Sequential()
      #input layer
      modell.add(Dense(26, input_dim = 26, activation = 'relu'))
      #hidden layer
      modell.add(Dense(50, activation = 'relu'))
      modell.add(Dense(25, activation = 'relu'))
      #output layer
      modell.add(Dense(1, kernel_initializer='normal'))

      # Step 2: Build the computational graph - compile
      modell.compile(loss = 'mean_absolute_error', optimizer = 'adam', metrics = ['mean_absolute_error'])

      # Step 3: Train the model
      modell.fit(X_train, Y_train, epochs = 30, batch_size = 50)
      scores = modell.evaluate(X_test, Y_test)
      print("\nks: %.2f%%" % (modell.metrics_names[1], scores[1]*100))

      from sklearn.metrics import r2_score, recall_score, precision_score

      Epoch 1/30
      28/28 [=====] - 0s 550us/step - loss: 213556.3906 - mean_absolute_error: 213
      556.3906
      Epoch 2/30
      28/28 [=====] - 0s 534us/step - loss: 213550.1250 - mean_absolute_error: 213
      550.1250
      Epoch 3/30
      28/28 [=====] - 0s 534us/step - loss: 213506.8906 - mean_absolute_error: 213
      506.8906
      Epoch 4/30
      28/28 [=====] - 0s 570us/step - loss: 213315.4375 - mean_absolute_error: 213
      315.4375
      Epoch 5/30
      28/28 [=====] - 0s 534us/step - loss: 212715.8750 - mean_absolute_error: 212
      715.8750
      Epoch 6/30
      28/28 [=====] - 0s 605us/step - loss: 211262.2656 - mean_absolute_error: 211
      262.2656
      Epoch 7/30
      28/28 [=====] - 0s 570us/step - loss: 208291.8594 - mean_absolute_error: 208
      291.8594
      Epoch 8/30
      28/28 [=====] - 0s 570us/step - loss: 202960.6562 - mean_absolute_error: 202
      960.6562
      Epoch 9/30
      28/28 [=====] - 0s 570us/step - loss: 194221.8594 - mean_absolute_error: 194
      221.8594
      Epoch 10/30
      28/28 [=====] - 0s 499us/step - loss: 180815.0469 - mean_absolute_error: 180
      815.0469
      Epoch 11/30
      28/28 [=====] - 0s 534us/step - loss: 161352.8750 - mean_absolute_error: 161
      352.8750
      Epoch 12/30
      28/28 [=====] - 0s 570us/step - loss: 134579.2188 - mean_absolute_error: 134
      579.2188
      Epoch 13/30
      28/28 [=====] - 0s 534us/step - loss: 101147.7422 - mean_absolute_error: 101
      147.7422
      Epoch 14/30
      28/28 [=====] - 0s 534us/step - loss: 71393.4375 - mean_absolute_error: 7139
      3.4375
      Epoch 15/30
      28/28 [=====] - 0s 570us/step - loss: 59262.0977 - mean_absolute_error: 5926
      2.0977
      Epoch 16/30
      28/28 [=====] - 0s 570us/step - loss: 57635.1797 - mean_absolute_error: 5763
      5.1797
      Epoch 17/30
      28/28 [=====] - 0s 534us/step - loss: 57333.2773 - mean_absolute_error: 5733
      3.2773
      Epoch 18/30
      28/28 [=====] - 0s 605us/step - loss: 57088.3594 - mean_absolute_error: 5708
      8.3594
      Epoch 19/30
      28/28 [=====] - 0s 570us/step - loss: 56872.6289 - mean_absolute_error: 5687
      2.6289
      Epoch 20/30
      28/28 [=====] - 0s 570us/step - loss: 56655.1914 - mean_absolute_error: 5665
      5.1914
      Epoch 21/30
      28/28 [=====] - 0s 499us/step - loss: 56440.2383 - mean_absolute_error: 5644
      0.2383
      Epoch 22/30
      28/28 [=====] - 0s 534us/step - loss: 56219.3320 - mean_absolute_error: 5621
      9.3320
      Epoch 23/30
      28/28 [=====] - 0s 534us/step - loss: 56037.3594 - mean_absolute_error: 5603
      7.3594
      Epoch 24/30
      28/28 [=====] - 0s 570us/step - loss: 55812.3086 - mean_absolute_error: 5581
      2.3086
      Epoch 25/30
      28/28 [=====] - 0s 570us/step - loss: 55634.8047 - mean_absolute_error: 5563
      4.8047
      Epoch 26/30
      28/28 [=====] - 0s 499us/step - loss: 55446.2578 - mean_absolute_error: 5544
      6.2578
      Epoch 27/30
      28/28 [=====] - 0s 534us/step - loss: 55287.0039 - mean_absolute_error: 5528
      7.0039
      Epoch 28/30
      28/28 [=====] - 0s 534us/step - loss: 55093.4844 - mean_absolute_error: 5509
      3.4844
      Epoch 29/30
      28/28 [=====] - 0s 534us/step - loss: 54943.8438 - mean_absolute_error: 5494
      3.8438
      Epoch 30/30
      28/28 [=====] - 0s 570us/step - loss: 54770.2305 - mean_absolute_error: 5477
      0.2383
      11/11 [=====] - 0s 545us/step - loss: 49416.7773 - mean_absolute_error: 4941
      6.7773

      mean_absolute_error: 4941677.73%
```

```
In [57]: Y_train_predict = modell.predict(X_train)
          Y_test_predict = modell.predict(X_test)

          print("Train score: {:.2f}".format(r2_score(Y_train, Y_train_predict)))
          print("Test score: {:.2f}".format(r2_score(Y_test, Y_test_predict)))

          Train score: 0.37
          Test score: 0.45
```

```
In [ ]:
```